MODUS Business Plan

Short-term plan (next 12 months) and 3 to 5 year plan

Prepared for Richard Latta, Post Oak Group. November 2026. Jeremiah Diana, Founder. jeremiah@moduspilot.com. moduspilot.com

All forward-looking figures in this document are projections. Each one is built from the assumptions printed next to it. None of them are results.

Part 1. Short-term plan (December 2026 to November 2027)

1. Where MODUS is today

MODUS is an answer engine that puts the leading AI models from several labs under one subscription. A user asks a question once. MODUS either routes it to the model best suited to the task or runs it on several models side by side and says which answer is strongest and why.

What is built and live:

- **Web app** at moduspilot.com.
- **Desktop apps** for Mac and Windows. The Mac app includes Screen Assist, which answers questions about whatever is on screen from a keyboard shortcut.
- **iPhone app** submitted to the App Store and in review.
- **Models from five labs.** The $28 plan includes GPT-6.1 Sol (OpenAI), Claude Sonnet 5.5 (Anthropic), Gemini 3.8 Flash and Gemini 3.5 Flash Lite (Google), DeepSeek V3.1 and Llama 3.3 (Meta). The PILOT plan adds Claude Opus 5.5, Claude Fable 5.1, GPT-6 Astra, Gemini 4 and Llama 4 Maverick.
- **Core features:** auto-routing, side-by-side compare, memory across conversations, connected inbox, calendar and files, actions that wait on user approval, projects, plugins and connectors to outside tools.
- **Billing** live through Stripe with usage limits that cap what any one account can cost us.

Pricing shown to customers:

Plan	Monthly	Annual	Includes
Free	$0	$0	Open models, rolling usage window
MODUS	$28	$280 ($23.33 per month)	Frontier models from five labs
PILOT	$59	$588 ($49 per month)	Top models, higher limits
Limits add-on	$10	not offered	More weekly headroom, stackable

What is not proven yet: that people who do not already know me will find MODUS, sign up and keep paying. The next 12 months are about proving exactly that.

The company is me, full time. I build the product and run distribution.

2. Product milestones

Quarter	Milestones
Q1 (Dec to Feb)	iOS approval and public App Store launch. Keep the model catalog current within two weeks of each major model release. Onboarding pass on web and mobile so a new user gets a useful answer in the first minute.
Q2 (Mar to May)	Android app built from the same codebase as iOS, submitted to Google Play. Referral credit inside the app. Usage reporting so I can see cost per user each month instead of estimating it.
Q3 (Jun to Aug)	Teams plan beta: shared billing, seats and shared projects for small companies. Windows desktop at feature parity with Mac, including Screen Assist.
Q4 (Sep to Nov)	Teams plan generally available. First localized pricing (two or three countries outside the US). Decision on a developer API based on demand seen during the year.

3. Customer acquisition by channel

Every channel is measured on the same funnel: visit, signup, first answer, paid. Signup and paid targets below are the base case from Part 2 (Section 3). They are targets, not forecasts I can promise.

Channel	What I will do	Target over 12 months
Agency creators and clients	I run an agency that works with creators and clients. Paid and affiliate placements with creators whose audiences already use AI tools.	8 to 12 creator placements per quarter. Largest single source of signups.
Paid ads (Meta first)	Small weekly tests on Instagram and Facebook. Kill any ad that does not produce signups at or under target cost within two weeks.	Cost per signup under $4 before scaling spend.
Affiliates	Commission on paid conversions through tracked codes. Recruit from creators, newsletter writers and power users.	25 active affiliates by Q4.
SEO and blog	Comparison pages and guides people search before choosing an AI tool. Two to four pieces a month.	15 percent of signups from search by Q4.
App Store launch	iOS launch in Q1, Android in Q2. Store listing tuned for search inside the stores.	Store installs tracked as their own channel from launch.
Reddit and communities	Posts and answers from my own account in AI and productivity communities. No fake accounts, no paid posting.	Steady referral traffic, tracked by link.

Quarterly signup and paying user targets are in Section 7.

4. Metrics I will report to Richard

A one-page report each month, same layout every time:

- Signups, by channel
- Activation: share of signups who get a full answer in their first session
- Free to paid conversion, by channel
- Paying users at month end, new and canceled
- Monthly churn of paying users
- Retention by monthly signup cohort at 30, 60 and 90 days
- AI cost per paying user and per free user
- Gross margin
- Cost per signup and cost per paying user, by paid channel

5. Use of funds and monthly burn

Spend is split by category. The split holds whatever the amount raised.

Category	Share	What it pays for
Distribution	28%	Creator placements, ad tests, affiliate payouts
Product and engineering	20%	Android build, contract help, design
Infrastructure	16%	Hosting, database, auth, email
Model usage	12%	AI provider bills for free and paid users
Founder draw	12%	Minimum to stay full time
Legal and accounting	7%	Entity, terms, app store and tax compliance
Reserve	5%	Held for a model price spike or a bad month

Monthly burn is kept flat and low on purpose. Distribution spend only goes up when a channel shows it can turn spend into paying users at a cost the margins below can carry.

6. Unit economics today (estimates)

Per paying user on the $28 plan, per month:

Line	Amount	Basis
Price	$28.00	Price on the site
Card processing	$1.11	Stripe standard rate, 2.9 percent plus $0.30
AI cost, typical	about $3.00	Heaviest week measured on a paying account, run for a full month
AI cost, hard ceiling	$18.90	The most the usage caps in our code allow one account to spend
Gross margin, typical	about 85%	($28.00 minus $1.11 minus $3.00) divided by $28.00
Gross margin, at the ceiling	about 29%	Worst case where a user hits every cap every week

The ceiling is enforced by the product, so even the heaviest possible user does not cost more than they pay. The free tier only offers open models with a rolling usage window. Its worst case is about $2 per active free user per month.

7. Quarterly milestones and targets

Quarter	Product	Distribution	Base case target
Q1	iOS live, onboarding pass	Creator placements start, first ad tests	1,200 signups, 58 paying
Q2	Android submitted, referrals	Affiliate program open, SEO cadence set	3,300 signups, 153 paying
Q3	Teams beta, Windows parity	Scale the best two channels	6,300 signups, 279 paying
Q4	Teams live, localized pricing	Cut channels that miss cost targets	10,200 signups, 435 paying

If Q1 and Q2 land well under these targets, I change channels before I change spend. The report in Section 4 is how Richard will see that happen.

Part 2. 3 to 5 year plan (2027 to 2031)

1. Market and competition

The market is large and people already use more than one assistant. Menlo Ventures estimates consumer AI spend reached $40 billion in 2026, up from $12 billion in 2025.[1] Their survey of 5,067 US adults found 55 percent of AI users now pay for at least one AI product, and the typical payer spends $20 to $49 a month.[1] The same study found the average user now uses 3.0 general AI assistants, up from 2.2 a year earlier.[1] Reach across ChatGPT (60 percent of AI users), Gemini (58 percent) and Claude (20 percent) adds up to well over 100 percent because people use more than one.[1]

That is the problem MODUS solves. People who want the best model for each task are paying for, or switching between, several subscriptions.

What the main options cost today:

Product	Entry paid plan	Top individual plan	Models from other labs
Claude (Anthropic)	Pro, $20 monthly or $17 a month annually[2]	Ultra, $200 a month[2]	No
Google AI (Gemini)	AI Plus, $4.99 a month. AI Pro, $19.99[3]	AI Ultra, from $99.99 to $199.99 a month[3]	No
ChatGPT (OpenAI)	Go and Plus plans[4]	Pro plan[4]	No
Perplexity	Pro	Max, with models from OpenAI, Anthropic and others in Research[5]	Yes
Poe (Quora)	Points-based plans	Points-based plans	Yes, one app across many models[6]
MODUS	$28 a month, $280 a year	PILOT, $59 a month	Yes, five labs

Where MODUS fits. The single-lab products are excellent at their own models and will keep getting better. They do not give you the other labs. Perplexity and Poe do offer several labs, so they are the closest competitors. Perplexity is built around search and research. Poe is built around a large catalog of bots. MODUS is built around getting one best answer: automatic routing, side-by-side compare with a verdict, memory and actions across your own email, calendar and files, on web, desktop and phone.

Competition I take seriously. Any of the large labs could add other labs' models, and Perplexity could push further into everyday assistant use. My response is in Section 6.

2. Product roadmap and expansion

Year	Expansion steps
2027 (Year 1)	iOS and Android live. Teams plan for small companies. Windows parity. First localized pricing.
2028 (Year 2)	Developer API: one key, models from every lab MODUS supports, billed by usage. Admin controls for teams (seat management, usage limits, data retention). More connectors (CRM, project tools, cloud storage). Five to eight localized markets.
2029 (Year 3)	Business tier with single sign-on, audit logs and data controls for larger companies. Education plan for students and schools. Native language support beyond English in the app.
2030 to 2031 (Years 4 and 5)	Enterprise sales motion. Partnerships that bundle MODUS with other software. Marketplace for plugins and workflows built by users and partners. Regional data hosting where customers require it.

Each step builds on what exists. Teams reuses the account and billing system. The API reuses the routing and cost controls that already run every chat. International reuses the same app with local prices and languages.

3. Revenue scenarios (projections)

These are projections. They cover the consumer and Teams subscriptions only. The API and enterprise lines in Section 2 are not modeled and would sit on top.

Assumptions used in every scenario:

- Blended price per paying user: **$30.63 a month.** Mix of 80 percent MODUS monthly at $28, 10 percent MODUS annual at $23.33 a month, 10 percent PILOT at $59.
- Card processing: **$1.19** per paying user per month, Stripe standard rate on that mix.
- Free tier cost: **$2 per signup** in their first month, the worst case for an active free user. Counted for every signup whether or not they are active.
- Signups grow in a straight line from the start to the end of each year.
- Paying users convert in the month they sign up. Churn applies to all paying users every month.
- Gross margin = revenue minus processing, AI cost and free tier cost, divided by revenue. Hosting, salaries and marketing are not in gross margin.

Assumptions that change by scenario:

Assumption	Conservative	Base	Upside
Signups per month, Year 1	200 rising to 600	300 rising to 1,400	500 rising to 2,500
Signups per month, Year 3	1,200 rising to 2,000	3,500 rising to 7,000	7,000 rising to 15,000
Signups per month, Year 5	3,000 rising to 4,000	11,000 rising to 16,000	25,000 rising to 40,000
Free to paid conversion	3%	5%	7%
Monthly churn, paying users	6%	4%	3%
AI cost per paying user per month	$6.00	$4.00	$3.00

The base case AI cost is above the $3 measured today, to allow for heavier users as the user base widens. The conservative case doubles it.

Results:

Scenario	Year	Signups in year	Paying users, year end	Revenue in year	Run rate at year end	Gross margin
Conservative	1	4,800	112	$19,900	$41,200	28%
Conservative	3	19,200	576	$161,500	$211,700	53%
Conservative	5	42,000	1,385	$432,800	$509,100	57%
Base	1	10,200	435	$70,800	$159,900	54%
Base	3	63,000	3,547	$921,600	$1,303,900	69%
Base	5	162,000	10,696	$3,203,600	$3,931,800	73%
Upside	1	18,000	1,119	$178,700	$411,300	66%
Upside	3	132,000	11,013	$2,761,900	$4,048,300	77%
Upside	5	390,000	38,948	$11,251,800	$14,317,100	79%

Run rate is paying users at year end times the blended price times 12. Margin is lower in early years because free tier cost per signup is large next to a small paying base. It climbs as paying users build up.

4. Team and hiring

Hiring follows revenue, not the other way around. Each hire below assumes the base case is on track.

Year	Team	First hires
2027	Founder full time, plus contractors	Contract mobile developer for Android, contract designer
2028	3 to 4 people	First full-time engineer, part-time customer support, growth and partnerships lead
2029	6 to 8 people	Second and third engineers (API and business features), full-time support, content and SEO
2030	10 to 14 people	Sales for teams and business accounts, security and compliance, data and finance
2031	15 to 20 people	Expansion by region and by product line as the numbers support it

5. Funding path

Stage	When	What has to be true
Current round	Now	Funds the 12-month plan in Part 1.
Pre-seed	After traction shows	Repeatable signups from at least two channels, conversion and churn near the base case for two quarters in a row.
Seed	Year 2 to 3	Teams plan selling, API in use, run rate on or above the base case.
Later rounds or profitability	Year 3 onward	Raise to grow faster only if margins hold. Otherwise grow from profit.

The business is set up so it does not need to raise to survive. Usage caps keep every paying account above water, and spend on distribution only rises with results.

6. Key risks and how I would handle them

Dependence on model provider APIs and pricing. MODUS runs on other companies' models. A provider could raise prices, cut access or change terms.

- Five labs today, so no one provider can switch MODUS off.
- Usage caps are set in cost, not message count, so a price rise shows up as lower limits on that model, not as a loss.
- When a provider raises prices, routing can send more everyday questions to cheaper models that answer them just as well.
- A reserve is held for a price spike while limits are adjusted.

Big players bundling multiple models. If a large lab or platform offers other labs' models in its own app, the main reason to choose MODUS gets weaker.

- Labs have little reason to promote a rival's models in their own product. A neutral product does not have that conflict.
- MODUS competes on the answer, not just access: routing, compare with a verdict, memory and actions on your own data.
- Teams, API and business tiers make MODUS harder to replace once a company has set it up.

Margins. Heavy users, expensive new models or a large free tier could squeeze gross margin.

- Every plan has a hard cost ceiling in code. The $28 plan cannot cost more than $18.90 in AI per user per month.
- I will track AI cost per user every month and report it (Part 1, Section 4).
- If margin falls below target, the levers are, in order: routing to cheaper models, tighter free tier limits, then pricing.

Distribution. The biggest near-term risk is not the product. It is getting it in front of enough of the right people at a cost the margins can carry. The 12-month plan tests six channels and cuts the ones that do not work, with the results reported monthly.

Founder dependence. Today the company is one person. The first full-time hire in 2028 is an engineer so the product does not depend on me alone. Code, accounts and operating docs are kept so someone else could pick them up.

Sources

1. Menlo Ventures, "2026: The State of Consumer AI," September 2026. https://menlovc.com/perspective/2026-the-state-of-consumer-ai/
2. Anthropic, Claude pricing page, accessed November 5, 2026. https://claude.com/pricing
3. Google, Google AI plans page (US), accessed November 5, 2026. https://gemini.google/subscriptions/
4. OpenAI, ChatGPT pricing page, plan names, accessed November 5, 2026. https://chatgpt.com/pricing
5. Perplexity Help Center, "Which Perplexity subscription plan is right for you," accessed November 5, 2026. https://www.perplexity.ai/help-center/en/articles/11187416-which-perplexity-subscription-plan-is-right-for-you
6. TechCrunch, "Poe's AI app now supports group chats across AI models," November 18, 2025. https://techcrunch.com/2025/11/18/poes-ai-app-now-supports-group-chats-across-ai-models/

MODUS cost figures (card processing, measured usage, cost ceilings) come from our own Stripe account and the usage limits in our product code. They are estimates where marked.